

R6 — extraction on 745: best variant A /
gemini:gemini-3.1-pro-preview = f1 0.091, uce 1.250 (31 rows ·
variants A,B,C,E · 5 models · projects
1729,3122,321,4006,745,OFR · levels extraction,offer)

Generated from canonical runs/leaderboard.csv (31 rows).

Gate: 745 extraction (group_f1≥0.7 AND unit_count_error≤0.05)

0/6 cells PASS, 6/6 FAIL.

model	group_f1	qty_mae	unit_count_error	cost_usd	elapsed_sec	gate
anthropic:claude-opus-4-8	0.000	—	0.708	\$2.029	154.8	FAIL (f1=0.000, uce=0.708; gate
gemini:gemini-3.1-pro-preview	0.091	5.000	1.250	\$0.202	170.1	FAIL (f1=0.091, uce=1.250; gate
anthropic:claude-opus-4-8	0.000	—	1.000	\$0.000	0.0	FAIL (f1=0.000, uce=1.000; gate
gemini:gemini-3.1-pro-preview	0.000	—	1.000	\$0.000	0.0	FAIL (f1=0.000, uce=1.000; gate
anthropic:claude-opus-4-8	0.056	0.000	1.833	\$1.902	255.9	FAIL (f1=0.056, uce=1.833; gate
gemini:gemini-3.1-pro-preview	0.000	—	1.417	\$0.309	339.3	FAIL (f1=0.000, uce=1.417; gate

Gate: 745 offer (group_f1≥0.7 AND unit_count_error≤0.05)

0/6 cells PASS, 6/6 FAIL.

model	group_f1	qty_mae	unit_count_error	cost_usd	elapsed_sec	gate
anthropic:claude-opus-4-8	0.000	—	0.708	\$2.029	154.8	FAIL (f1=0.000, uce=0.708; gate
gemini:gemini-3.1-pro-preview	0.100	2.000	1.250	\$0.202	170.1	FAIL (f1=0.100, uce=1.250; gate
anthropic:claude-opus-4-8	0.000	—	1.000	\$0.000	0.0	FAIL (f1=0.000, uce=1.000; gate
gemini:gemini-3.1-pro-preview	0.000	—	1.000	\$0.000	0.0	FAIL (f1=0.000, uce=1.000; gate
anthropic:claude-opus-4-8	0.000	—	1.833	\$1.902	255.9	FAIL (f1=0.000, uce=1.833; gate
gemini:gemini-3.1-pro-preview	0.080	0.000	1.417	\$0.309	339.3	FAIL (f1=0.080, uce=1.417; gate

Leaderboard (variant × model × project × gt_level)									
	project	gt_level	group_f1	qty_mae	unit_count_error	glass_acc	cost_usd	elapsed_sec	error
ot-4.1-nano	1729	—	—	—	—	—	—	—	—
ot-4o	1729	—	—	—	—	—	—	—	—
gemini-3.1-pro-preview	1729	—	—	—	—	—	—	—	—
ot-4.1-nano	1729	—	—	—	—	—	—	—	—
ot-4.1-nano	3122	—	—	—	—	—	—	—	—
ot-4o	3122	—	—	—	—	—	—	—	—
gemini-3.1-pro-preview	3122	—	—	—	—	—	—	—	—
ot-4.1-nano	321	—	—	—	—	—	—	—	—
gemini-3.1-pro-preview	321	—	—	—	—	—	—	—	—
ot-4.1-nano	4006	—	0.000	0.000	0.330	1.000	—	—	—
ot-4o	4006	—	0.000	0.000	0.687	1.000	—	—	—
ot-4.1-nano	745	—	—	—	—	—	—	—	—
ot-4o	745	—	—	—	—	—	—	—	—
gemini-3.1-pro-preview	745	—	—	—	—	—	—	—	—
ot-4.1-nano	745	—	—	—	—	—	—	—	—
c:claude-opus-4-8	745	extraction	0.000	—	0.708	—	\$2.029	154.8	—
gemini-3.1-pro-preview	745	extraction	0.091	5.000	1.250	—	\$0.202	170.1	—
c:claude-opus-4-8	745	extraction	0.000	—	1.000	—	\$0.000	0.0	—
gemini-3.1-pro-preview	745	extraction	0.000	—	1.000	—	\$0.000	0.0	—
c:claude-opus-4-8	745	extraction	0.056	0.000	1.833	—	\$1.902	255.9	—
gemini-3.1-pro-preview	745	extraction	0.000	—	1.417	—	\$0.309	339.3	—
c:claude-opus-4-8	745	offer	0.000	—	0.708	—	\$2.029	154.8	—
gemini-3.1-pro-preview	745	offer	0.100	2.000	1.250	—	\$0.202	170.1	—
c:claude-opus-4-8	745	offer	0.000	—	1.000	—	\$0.000	0.0	—
gemini-3.1-pro-preview	745	offer	0.000	—	1.000	—	\$0.000	0.0	—
c:claude-opus-4-8	745	offer	0.000	—	1.833	—	\$1.902	255.9	—
gemini-3.1-pro-preview	745	offer	0.080	0.000	1.417	—	\$0.309	339.3	—
c:claude-opus-4-7	OFR	—	0.533	6.875	0.000	1.000	—	—	—
ot-4.1-nano	OFR	—	—	—	—	—	—	—	—
ot-4o	OFR	—	0.581	2.000	0.226	1.000	—	—	—
gemini-3.1-pro-preview	OFR	—	0.714	8.000	0.000	1.000	—	—	—

Variant summary (A-F) — extraction-gate pass rate on 745

variant	cells	PASS	FAIL	PASS %	best group_f1	min cost_usd
A	2	0	2	0%	0.091	\$0.202
B	0	—	—	—	—	—
C	2	0	2	0%	0.000	\$0.000
D	0	—	—	—	—	—
E	2	0	2	0%	0.056	\$0.309
F	0	—	—	—	—	—

Delta vs R5

variant	model	project	R6 f1	R5 f1	$\Delta f1$	R6 uce	R5 uce	Δuce
E	gemini:gemini-3.1-pro-preview	745	0.000	0.000	+0.000	1.417	0.000	+1.417
A	gemini:gemini-3.1-pro-preview	745	0.091	0.000	+0.091	1.250	0.000	+1.250
C	gemini:gemini-3.1-pro-preview	745	0.000	0.000	+0.000	1.000	0.000	+1.000

What this report does NOT claim

This report does NOT claim:

- that the extraction-gate has been taken on 745. Where the gate fails it is reported with the literal f1/uce numbers; we did not lower thresholds to make a cell pass.
- that ground truth was edited to match any model output. The *745_extract.json* file was rewritten once, before any benchmark, from a human-verified dual-VLM transcription of sheet A1.0 (R6 step 1.4 confirmed by V on 2026-06-08).
- that the offer GT (*tests/ground_truth/745.json*) is correct. It pre-dates R6 verification and uses fabricated CM1-4 / DR1-2 mark IDs. The offer gate numbers here are an honest measurement of where the discovery pipeline ends up vs that stale offer GT, NOT evidence that the discovery layer is broken.
- that Variant C was budget-tested on 745. It was parked by the R6 cost gate (see *docs/r6_variant_c_cost_gate.md*); rows for C on 745 carry the literal *all_pages_raster* error string instead of a fabricated metric.
- that the "dash" (—) cells in any table mean zero. They mean the metric was not computable on that cell (no matched groups, no GT, or the provider errored before metrics ran).